D07 | XOR decision block

All-case aligned observations: 21,588 | Train: 15,018 | Held-out test: 6,570

Test accuracy 0.877 | Majority baseline 0.826 | Balanced accuracy 0.904 | Macro-F1 0.824

B1 -> O_Accepted | aligned support 17,873
B2 -> A_Denied | aligned support 3,715

Learned rules (training cases only; depth <= 3; minimum leaf size 50)

```
e.CreditScore <= 296 AND c.AppType != "New credit" AND e.NumberOfTerms <= 104.5 -> B1 [train n=622]

e.CreditScore <= 296 AND c.AppType != "New credit" AND e.NumberOfTerms > 104.5 -> B1 [train n=965]

e.CreditScore <= 296 AND c.AppType = "New credit" AND e.OfferedAmount <= 13450 -> B2 [train n=1,921]

e.CreditScore <= 296 AND c.AppType = "New credit" AND e.OfferedAmount > 13450 -> B2 [train n=2,209]

e.CreditScore > 296 -> B1 [train n=9,301]
```

Branches never predicted by this tree: none

c. = case input at start; e. = last synchronized event input. Missing numeric value = -1; explicit *_known flags.
Inputs update after synchronous moves only. One optimal alignment per variant; alternative alignments may change decisions.
Statistical associations; not bank policy. Structural replay coverage 80.45% is measured without these guards.